

# data-exploration-and-cleaning

February 18, 2022

## 1 Data Cleaning and Exploration Notebook

In this notebook, I'm going to import data collected from various sources, clean it, make it uniform and merge it together before exporting it to excel as a dataset containing everything I'll use for my analysis.

### 1.1 Contents

1. Setup
2. Import Datasets
  - Convert Data from Wide to Long
  - Inspect Data
3. Adjust for Population and Inflation and Explore Data
  - Prices
  - Tariffs and Imports
  - U.S. Production
  - Exports
4. Combine Datasets

### 1.2 Setup

```
[1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import datetime as dt
import calendar
import math
import re
```

```
[2]: # Create a map of 3-letter month to 2-digit month
month_map = {'JAN': '01', 'FEB': '02', 'MAR': '03', 'APR': '04', 'MAY': '05', 'JUN':
↪ '06', 'JUL': '07', 'AUG': '08', 'SEP': '09', 'OCT': '10', 'NOV': '11', 'DEC': '12'}
```

```
[3]: # Pathnames for various datasets
avg_price_data = '../data/avg-price-us-cities.xlsx'
tariff_active_data = '../data/ukfrspger_25.xlsx'
ukfrspde_imports_data = '../data/usitc-ukfrspde-imports.xlsx'
world_imports_data = '../data/usitc-world-wine-imports.xlsx'
```

```
us_production_data = '../data/wine-prod-ttb.xlsx'
us_exports_data = '../data/exports-monetary-value-and-liters.xlsx'
ppi_data = '../data/PCU3121303121300.xls'
population_data = '../data/POPTHM.xls'
```

```
[4]: # Making plots a bit more accessible
sns.set_palette('colorblind')
```

### 1.3 Import Datasets

```
[5]: production_df = pd.read_excel(io=us_production_data, header=3)
exports_dict = pd.read_excel(io=us_exports_data, sheet_name=['FAS Value', 'First_
↳Unit of Quantity'])
wine_prices_df = pd.read_excel(io=avg_price_data)
imports_world_dict = pd.read_excel(world_imports_data, sheet_name=None)
imports_ukfrspde_dict = pd.read_excel(ukfrspde_imports_data, sheet_name=None)
tariff_df = pd.read_excel(tariff_active_data)
ppi_df = pd.read_excel(io=ppi_data, header=10)
population_df = pd.read_excel(io=population_data, header=10)
```

#### 1.3.1 Convert Data Presentation to Long

The exports and imports datasets are in a wide format where each row corresponds to a year and each column a month of that year. The sheets of these datasets are different datapoints being reported.

I'll want to convert the rows to correspond to months in the format YYYY-MM and the columns to the specific datapoints being reported.

Let's define a function to help convert the datasets.

```
[6]: def wide_to_long(df, col_name):
    df = pd.melt(df, id_vars=['Year'], var_name='month', value_name=col_name)

    df['month'] = df['month'].map(month_map)
    df['month'] = df['Year'].astype(str) + '-' + df['month']

    df.drop(columns='Year', inplace=True)
    return df
```

Let's add another function to get the last day of the month.

```
[7]: def last_day_of_month(month):
    if pd.isnull(month):
        return

    d = month.split('-')
    date = dt.date(pd.to_numeric(d[0]), pd.to_numeric(d[1]), 1)
    return date.replace(day = calendar.monthrange(date.year, date.month)[1])
```

We'll also need a time series for joining the data together.

```
[8]: months = pd.date_range('2000-01-01', '2022-01-01', freq='M')
```

Now we can get the imports data as a long dataset.

We'll first grab the imports data for the entire world and then we'll go and get the data for just the U.K., France, Spain, and Germany.

```
[9]: imports_world_df = pd.DataFrame()
imports_world_df['month'] = months

for name, sheet in imports_world_dict.items():
    if name != 'Query Parameters' and name != 'Query Results':
        sheet_df = pd.DataFrame(sheet)
        sheet_df.columns = sheet_df.iloc[0]
        sheet_df = sheet_df.iloc[1: , :]
        sheet_df = sheet_df[['Year', 'JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN',
        → 'JUL', 'AUG', 'SEP', 'OCT', 'NOV', 'DEC']]

        temp_df = wide_to_long(sheet_df, col_name=name)
        temp_df['month'] = temp_df['month'].map(lambda x: last_day_of_month(x))
        temp_df['month'] = temp_df['month'].astype('datetime64[ns]')
        imports_world_df = imports_world_df.merge(temp_df, on='month')

imports_world_df.sort_values('month', ascending=True, inplace=True)
imports_world_df.head()
```

```
[9]:      month Dutiable Value Landed Duty-Paid Value Customs Value \
0 2000-01-31      133791586              144253063      134876329
1 2000-02-29      140504325              149395178      141089327
2 2000-03-31      168440973              181115109      170803664
3 2000-04-30      174444956              186771728      175506155
4 2000-05-31      177795372              190281882      178396150

      First Unit of Quantity Charges, Insurance, and Freight Calculated Duties
0              30869692              7058649      2318085
1              27334547              6156122      2149729
2              33950034              7719881      2591564
3              37381318              8447127      2818446
4              38726698              9091316      2794416
```

Alright, let's pull the data for the countries that the U.S. levied the higher tariffs on.

```
[10]: imports_ukfrspde_dict = pd.read_excel(ukfrspde_imports_data, sheet_name=None)

imports_ukfrspde_df = pd.DataFrame()
imports_ukfrspde_df['month'] = months
imports_uk_df = imports_ukfrspde_df.copy()
```

```

imports_fr_df = imports_ukfrspde_df.copy()
imports_de_df = imports_ukfrspde_df.copy()
imports_sp_df = imports_ukfrspde_df.copy()

for name, sheet in imports_ukfrspde_dict.items():
    if name != 'Query Parameters' and name != 'Query Results':
        sheet_df = pd.DataFrame(sheet)
        sheet_df.columns = sheet_df.iloc[0]
        sheet_df = sheet_df.iloc[1: , :]
        sheet_df = sheet_df[['Country', 'Year', 'JAN', 'FEB', 'MAR', 'APR',
↪ 'MAY', 'JUN', 'JUL', 'AUG', 'SEP', 'OCT', 'NOV', 'DEC']]

        temp_fr_df = sheet_df.loc[(sheet_df['Country'] == 'France')]
        temp_de_df = sheet_df.loc[(sheet_df['Country'] == 'Germany')]
        temp_sp_df = sheet_df.loc[(sheet_df['Country'] == 'Spain')]
        temp_uk_df = sheet_df.loc[(sheet_df['Country'] == 'United Kingdom')]
        temp_all_df = sheet_df.loc[(sheet_df['Country'].str.contains('Total'))]

        temp_fr_df2 = wide_to_long(temp_fr_df, col_name=name)
        temp_de_df2 = wide_to_long(temp_de_df, col_name=name)
        temp_sp_df2 = wide_to_long(temp_sp_df, col_name=name)
        temp_uk_df2 = wide_to_long(temp_uk_df, col_name=name)
        temp_all_df2 = wide_to_long(temp_all_df, col_name=name)

        temp_fr_df2['month'] = temp_fr_df2['month'].map(lambda x:
↪ last_day_of_month(x))
        temp_fr_df2['month'] = temp_fr_df2['month'].astype('datetime64[ns]')
        imports_fr_df = imports_fr_df.merge(temp_fr_df2, on='month')

        temp_de_df2['month'] = temp_de_df2['month'].map(lambda x:
↪ last_day_of_month(x))
        temp_de_df2['month'] = temp_de_df2['month'].astype('datetime64[ns]')
        imports_de_df = imports_de_df.merge(temp_de_df2, on='month')

        temp_sp_df2['month'] = temp_sp_df2['month'].map(lambda x:
↪ last_day_of_month(x))
        temp_sp_df2['month'] = temp_sp_df2['month'].astype('datetime64[ns]')
        imports_sp_df = imports_sp_df.merge(temp_sp_df2, on='month')

        temp_uk_df2['month'] = temp_uk_df2['month'].map(lambda x:
↪ last_day_of_month(x))
        temp_uk_df2['month'] = temp_uk_df2['month'].astype('datetime64[ns]')
        imports_uk_df = imports_uk_df.merge(temp_uk_df2, on='month')

        temp_all_df2['month'] = temp_all_df2['month'].map(lambda x:
↪ last_day_of_month(x))

```

```

temp_all_df2['month'] = temp_all_df2['month'].astype('datetime64[ns]')
imports_ukfrspde_df = imports_ukfrspde_df.merge(temp_all_df2,
↳ on='month')

imports_ukfrspde_df.sort_values('month', ascending=True, inplace=True)
imports_ukfrspde_df.head()

```

```

[10]:      month Charges, Insurance, and Freight Calculated Duties Dutiable Value \
0 2000-01-31                2406931                699303                56667673
1 2000-02-29                2356486                857890                73857433
2 2000-03-31                2804317                914008                83440008
3 2000-04-30                2989501                974797                83020862
4 2000-05-31                3037785                897331                75535353

      Landed Duty-Paid Value First Unit of Quantity Customs Value
0                59812261                8768676                56706027
1                77087577                8961916                73873201
2                87219165                10474993                83500840
3                87040067                11128077                83075769
4                79534639                10874051                75599523

```

Now let's make the exports data long.

```

[11]: for name, sheet in exports_dict.items():
    if name == 'FAS Value':
        exports_value_df = pd.DataFrame(sheet)
        exports_value_df.columns = exports_value_df.iloc[0]
        exports_value_df = exports_value_df.iloc[1: , :]
        exports_value_df = wide_to_long(exports_value_df, col_name=name)
    if name == 'First Unit of Quantity':
        exports_quantity_df = pd.DataFrame(sheet)
        exports_quantity_df.columns = exports_quantity_df.iloc[0]
        exports_quantity_df = exports_quantity_df.iloc[1: , :]
        exports_quantity_df = wide_to_long(exports_quantity_df, col_name=name)

exports_df = exports_value_df.merge(exports_quantity_df, on='month')

# There are 108 empty rows at the bottom of the dataframe due to extra cells in
↳ the
# excel spreadsheet. After checking, these are the only empty rows and dropping
↳ them
# leaves 264 entries, as there should be.
exports_df.dropna(inplace=True)

exports_df['month'] = exports_df['month'].map(lambda x: last_day_of_month(x))
exports_df.sort_values('month', ascending=True, inplace=True)

```

```
exports_df.head()
```

```
[11]:      month FAS Value First Unit of Quantity
0    2000-01-31  45525687          22385348
31   2000-02-29  41336082          20954045
62   2000-03-31  41446856          20240059
93   2000-04-30  41411758          21745478
124  2000-05-31  42672115          22059299
```

### 1.3.2 Inspect Data

Here I'll provide summary statistics for each of the datasets before we go on to do any transformations on the data.

**Producer Price Index** I'm going to adjust for inflation using the industry-specific series for the producer price index as provided by the Federal Reserve (FRED). I'm doing this because most of the data used is production-side and this provides a more-specific inflation adjustment for this industry than the regular CPI.

The series is PCU3121303121300.

FRED summarizes the dataset as follows: Producer Price Index by Industry: Wineries: Wines, Brandy, and Brandy Spirits, Index Dec 1998=100, Monthly, Not Seasonally Adjusted

```
[12]: # Get the last day of the month
ppi_df['month'] = ppi_df['observation_date'].map(lambda x: str(x.year) + '-' +
    ↪ '{:02d}'.format(x.month))
ppi_df['month'] = ppi_df['month'].map(lambda x: last_day_of_month(x))
ppi_df['month'] = ppi_df['month'].astype('datetime64[ns]')
```

```
[13]: display(ppi_df.head())
display(ppi_df.info())
display(ppi_df.describe())
```

	observation_date	PCU3121303121300	month
0	2000-01-01	103.0	2000-01-31
1	2000-02-01	101.1	2000-02-29
2	2000-03-01	103.0	2000-03-31
3	2000-04-01	101.0	2000-04-30
4	2000-05-01	101.9	2000-05-31

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 265 entries, 0 to 264
Data columns (total 3 columns):
#   Column                Non-Null Count  Dtype
---  -
0   observation_date      265 non-null   datetime64[ns]
1   PCU3121303121300     265 non-null   float64
2   month                 265 non-null   datetime64[ns]
```

```
dtypes: datetime64[ns](2), float64(1)
memory usage: 6.3 KB
```

None

```
PCU3121303121300
count      265.000000
mean       116.637834
std         9.501049
min        100.900000
25%        106.000000
50%        118.800000
75%        123.400000
max        134.442000
```

**U.S. Population** Population Data provided by the Federal Reserve (FRED). I'll be using this to adjust for per-capita wine quantities. This data is presently in thousands.

POPTHM: Population, Thousands, Monthly, Not Seasonally Adjusted

```
[14]: population_df['month'] = population_df['observation_date'].map(lambda x: str(x.
    ↪year) + '-' + '{:02d}'.format(x.month))
population_df['month'] = population_df['month'].map(lambda x:
    ↪last_day_of_month(x))
population_df['month'] = population_df['month'].astype('datetime64[ns]')
population_df['population'] = population_df['POPTHM'] * 1000
```

```
[15]: display(population_df.info())
display(population_df.describe())
display(population_df.head())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 264 entries, 0 to 263
Data columns (total 4 columns):
#   Column          Non-Null Count  Dtype
---  -
0   observation_date  264 non-null   datetime64[ns]
1   POPTHM           264 non-null   int64
2   month            264 non-null   datetime64[ns]
3   population        264 non-null   int64
dtypes: datetime64[ns](2), int64(2)
memory usage: 8.4 KB
```

None


```
POPTHM    population
count      264.000000  2.640000e+02
mean    309321.689394  3.093217e+08
std      15255.933065  1.525593e+07
min      281083.000000  2.810830e+08
```

25%	296013.750000	2.960138e+08
50%	310940.500000	3.109405e+08
75%	323112.000000	3.231120e+08
max	331895.000000	3.318950e+08

	observation_date	POPTHM	month	population
0	2000-01-01	281083	2000-01-31	281083000
1	2000-02-01	281299	2000-02-29	281299000
2	2000-03-01	281531	2000-03-31	281531000
3	2000-04-01	281763	2000-04-30	281763000
4	2000-05-01	281996	2000-05-31	281996000

```
[16]: sns.lineplot(data=population_df, x='month', y='POPTHM')
```

```
[16]: <AxesSubplot:xlabel='month', ylabel='POPTHM'>
```



## World Imports

```
[17]: display(imports_world_df.head())
display(imports_world_df.describe())
display(imports_world_df.info())
```

	month	Dutiable Value	Landed Duty-Paid Value	Customs Value \
0	2000-01-31	133791586	144253063	134876329
1	2000-02-29	140504325	149395178	141089327
2	2000-03-31	168440973	181115109	170803664
3	2000-04-30	174444956	186771728	175506155



```
4 2000-05-31      177795372      190281882      178396150
```

```
First Unit of Quantity Charges, Insurance, and Freight Calculated Duties
0      30869692      7058649      2318085
1      27334547      6156122      2149729
2      33950034      7719881      2591564
3      37381318      8447127      2818446
4      38726698      9091316      2794416
```

```
<ipython-input-17-09d884bef029>:2: FutureWarning: Treating datetime data as
categorical rather than numeric in `.describe` is deprecated and will be removed
in a future version of pandas. Specify `datetime_is_numeric=True` to silence
this warning and adopt the future behavior now.
```

```
display(imports_world_df.describe())
```

```

count      month  Dutiable Value  Landed Duty-Paid Value \
unique      264      264.0      264.0
top  2009-05-31 00:00:00  387436035.0      389987329.0
freq      1      1.0      1.0
first  2000-01-31 00:00:00      NaN      NaN
last   2021-12-31 00:00:00      NaN      NaN
```

```

Customs Value  First Unit of Quantity \
count      264.0      264.0
unique      264.0      264.0
top   305807449.0      67360257.0
freq      1.0      1.0
first      NaN      NaN
last      NaN      NaN
```

```

Charges, Insurance, and Freight  Calculated Duties
count      264.0      264.0
unique      264.0      264.0
top      12248065.0      5050881.0
freq      1.0      1.0
first      NaN      NaN
last      NaN      NaN
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 264 entries, 0 to 263
```

```
Data columns (total 7 columns):
```

```

#   Column      Non-Null Count  Dtype
---  -
0   month      264 non-null    datetime64[ns]
1   Dutiable Value      264 non-null    object
2   Landed Duty-Paid Value      264 non-null    object
3   Customs Value      264 non-null    object
4   First Unit of Quantity      264 non-null    object
```

```

5   Charges, Insurance, and Freight  264 non-null    object
6   Calculated Duties                264 non-null    object
dtypes: datetime64[ns](1), object(6)
memory usage: 16.5+ KB

```

None

I'll clean up some of these variable names to make them cleaner.

```

[18]: def reformat_column_names(col_name):
        return str.lower(re.subn(' | |-', '_', col_name)[0])

```

Before looking into the values' distributions, I'm going to adjust for inflation and population.

```

[19]: imports_world_df = imports_world_df.merge(ppi_df, on='month')
imports_world_df = imports_world_df.merge(population_df[['month',
↳ 'population']], on='month')

for c in imports_world_df.columns:
    col_name = reformat_column_names(c) + '_adj'
    if c == 'population':
        imports_world_df['wine_quantity_per_capita'] = imports_world_df['First_
↳ Unit of Quantity'] / imports_world_df[c].astype('f4')
    elif c != 'First Unit of Quantity' and c != 'month' and c !=
↳ 'PCU3121303121300' and c != 'observation_date':
        imports_world_df[col_name] = imports_world_df[c] /
↳ imports_world_df['PCU3121303121300'] * 100

imports_world_df.head()

```

```

[19]:      month Dutiable Value Landed Duty-Paid Value Customs Value \
0 2000-01-31      133791586              144253063      134876329
1 2000-02-29      140504325              149395178      141089327
2 2000-03-31      168440973              181115109      170803664
3 2000-04-30      174444956              186771728      175506155
4 2000-05-31      177795372              190281882      178396150

      First Unit of Quantity Charges, Insurance, and Freight Calculated Duties \
0              30869692              7058649              2318085
1              27334547              6156122              2149729
2              33950034              7719881              2591564
3              37381318              8447127              2818446
4              38726698              9091316              2794416

      observation_date PCU3121303121300 population dutiable_value_adj \
0      2000-01-01              103.0      281083000      129894743.68932
1      2000-02-01              101.1      281299000      138975593.47181
2      2000-03-01              103.0      281531000      163534925.242718
3      2000-04-01              101.0      281763000      172717778.217822

```

```
4      2000-05-01      101.9    281996000    174480247.301276
```

```

landed_duty_paid_value_adj customs_value_adj \
0      140051517.475728    130947892.23301
1      147769711.177052    139554230.464886
2      175839911.650485      165828800.0
3      184922502.970297    173768470.29703
4      186733937.193327    175069823.356232

```

```

charges_insurance_and_freight_adj calculated_duties_adj \
0      6853057.281553      2250567.961165
1      6089141.444115      2126339.268051
2      7495030.097087      2516081.553398
3      8363492.079208      2790540.594059
4      8921801.766438      2742312.070658

```


```

wine_quantity_per_capita
0      0.109824
1      0.097173
2      0.120591
3      0.132669
4      0.137331

```

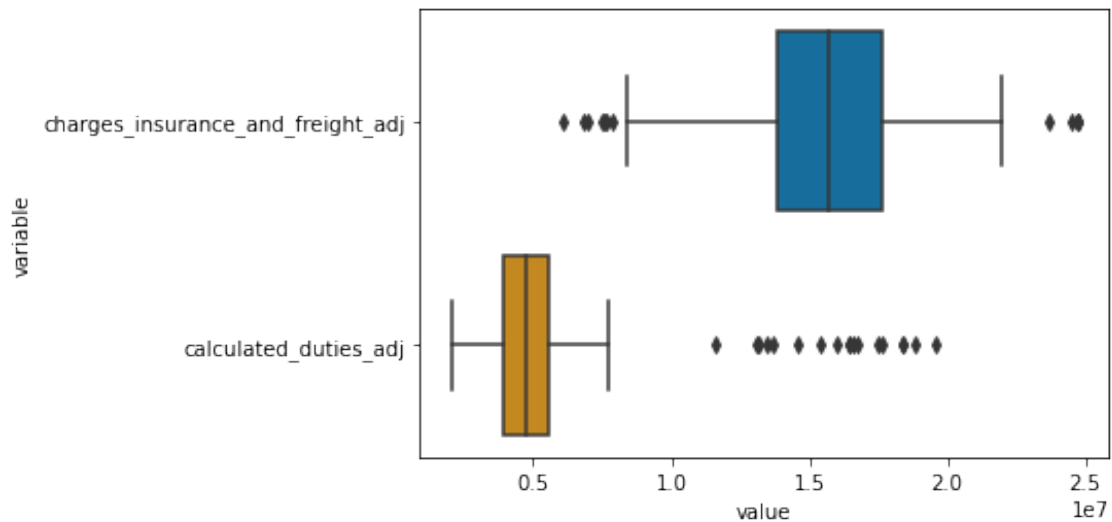
```
[20]: imports_value_boxplot_data = pd.melt(imports_world_df[['dutiabile_value_adj',
↳ 'landed_duty_paid_value_adj', 'customs_value_adj']])
sns.boxplot(data=imports_value_boxplot_data, y='variable', x='value')
```

```
[20]: <AxesSubplot:xlabel='value', ylabel='variable'>
```




```
[21]: imports_boxplot_data = pd.
      ↪ melt(imports_world_df[['charges_insurance_and_freight_adj',
      ↪ 'calculated_duties_adj']])
      sns.boxplot(data=imports_boxplot_data, y='variable', x='value')
```

[21]: <AxesSubplot:xlabel='value', ylabel='variable'>




```
[22]: sns.boxplot(data=imports_world_df['First Unit of Quantity'])
```

[22]: <AxesSubplot:>



```
[23]: sns.boxplot(data=imports_world_df['wine_quantity_per_capita'])
```

```
[23]: <AxesSubplot:>
```



## U.K., France, Germany, Spain Imports

```
[24]: imports_ukfrspde_df = imports_ukfrspde_df.merge(ppi_df, on='month')
imports_ukfrspde_df = imports_ukfrspde_df.merge(population_df[['month',
↳ 'population']], on='month')

for c in imports_ukfrspde_df.columns:
    if c == 'population':
        imports_ukfrspde_df['wine_quantity_per_capita'] =
↳ imports_ukfrspde_df['First Unit of Quantity'] /
↳ imports_ukfrspde_df['population']
    elif c != 'First Unit of Quantity' and c != 'month' and c !=
↳ 'PCU3121303121300' and c != 'observation_date':
        col_name = reformat_column_names(c) + '_adj'
        imports_ukfrspde_df[col_name] = imports_ukfrspde_df[c] /
↳ imports_ukfrspde_df['PCU3121303121300'] * 100

imports_ukfrspde_df.head()
```

```
[24]:      month Charges, Insurance, and Freight Calculated Duties Dutiable Value \
0 2000-01-31      2406931      699303      56667673
1 2000-02-29      2356486      857890      73857433
2 2000-03-31      2804317      914008      83440008
3 2000-04-30      2989501      974797      83020862
4 2000-05-31      3037785      897331      75535353
```

```
      Landed Duty-Paid Value First Unit of Quantity Customs Value \
0      59812261      8768676      56706027
1      77087577      8961916      73873201
2      87219165      10474993      83500840
3      87040067      11128077      83075769
4      79534639      10874051      75599523
```

```
      observation_date PCU3121303121300 population \
0      2000-01-01      103.0      281083000
1      2000-02-01      101.1      281299000
2      2000-03-01      103.0      281531000
3      2000-04-01      101.0      281763000
4      2000-05-01      101.9      281996000
```

```
      charges_insurance_and_freight_adj calculated_duties_adj dutiable_value_adj \
0      2336826.213592      678934.951456      55017158.252427
1      2330846.686449      848555.885262      73053840.751731
2      2722637.864078      887386.407767      81009716.504854
3      2959901.980198      965145.544554      82198873.267327
4      2981143.277723      880599.607458      74126941.118744
```

```
      landed_duty_paid_value_adj customs_value_adj wine_quantity_per_capita
```

0	58070156.31068	55054395.145631	0.031196
1	76248839.762611	73069437.1909	0.031859
2	84678800.970874	81068776.699029	0.037207
3	86178284.158416	82253236.633663	0.039494
4	78051657.50736	74189914.622179	0.038561

```
[25]: display(imports_ukfrspde_df.head())
display(imports_ukfrspde_df.describe())
display(imports_ukfrspde_df.info())
```

	month	Charges, Insurance, and Freight	Calculated Duties	Dutiable Value \
0	2000-01-31	2406931	699303	56667673
1	2000-02-29	2356486	857890	73857433
2	2000-03-31	2804317	914008	83440008
3	2000-04-30	2989501	974797	83020862
4	2000-05-31	3037785	897331	75535353

	Landed Duty-Paid Value	First Unit of Quantity	Customs Value \
0	59812261	8768676	56706027
1	77087577	8961916	73873201
2	87219165	10474993	83500840
3	87040067	11128077	83075769
4	79534639	10874051	75599523

	observation_date	PCU3121303121300	population \
0	2000-01-01	103.0	281083000
1	2000-02-01	101.1	281299000
2	2000-03-01	103.0	281531000
3	2000-04-01	101.0	281763000
4	2000-05-01	101.9	281996000

	charges_insurance_and_freight_adj	calculated_duties_adj	dutiable_value_adj \
0	2336826.213592	678934.951456	55017158.252427
1	2330846.686449	848555.885262	73053840.751731
2	2722637.864078	887386.407767	81009716.504854
3	2959901.980198	965145.544554	82198873.267327
4	2981143.277723	880599.607458	74126941.118744

	landed_duty_paid_value_adj	customs_value_adj	wine_quantity_per_capita
0	58070156.31068	55054395.145631	0.031196
1	76248839.762611	73069437.1909	0.031859
2	84678800.970874	81068776.699029	0.037207
3	86178284.158416	82253236.633663	0.039494
4	78051657.50736	74189914.622179	0.038561

	PCU3121303121300	population
count	264.000000	2.640000e+02
mean	116.571076	3.093217e+08

```
std          9.456620  1.525593e+07
min          100.900000  2.810830e+08
25%          105.925000  2.960138e+08
50%          118.800000  3.109405e+08
75%          123.175000  3.231120e+08
max          134.442000  3.318950e+08
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 264 entries, 0 to 263
Data columns (total 16 columns):
```

#	Column	Non-Null Count	Dtype
0	month	264 non-null	datetime64[ns]
1	Charges, Insurance, and Freight	264 non-null	object
2	Calculated Duties	264 non-null	object
3	Dutiable Value	264 non-null	object
4	Landed Duty-Paid Value	264 non-null	object
5	First Unit of Quantity	264 non-null	object
6	Customs Value	264 non-null	object
7	observation_date	264 non-null	datetime64[ns]
8	PCU3121303121300	264 non-null	float64
9	population	264 non-null	int64
10	charges_insurance_and_freight_adj	264 non-null	object
11	calculated_duties_adj	264 non-null	object
12	dutiable_value_adj	264 non-null	object
13	landed_duty_paid_value_adj	264 non-null	object
14	customs_value_adj	264 non-null	object
15	wine_quantity_per_capita	264 non-null	object

```
dtypes: datetime64[ns](2), float64(1), int64(1), object(12)
```


```
memory usage: 35.1+ KB
```

```
None
```

```
[26]: imports_value_boxplot_ukfrspde_df = pd.
      ↪ melt(imports_ukfrspde_df[['dutiable_value_adj',
      ↪ 'landed_duty_paid_value_adj', 'customs_value_adj']])
      sns.boxplot(data=imports_value_boxplot_ukfrspde_df, y='variable', x='value')
```

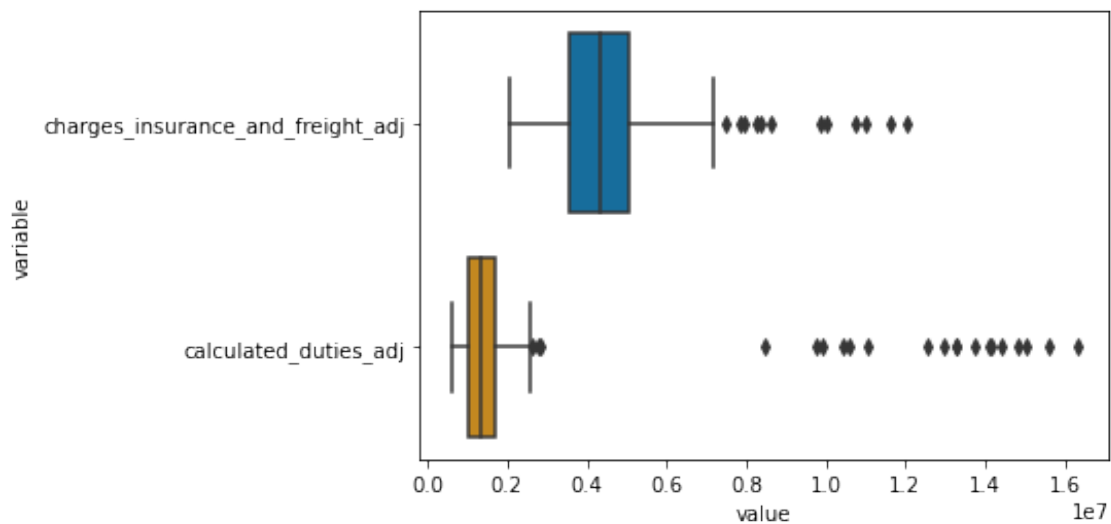
```
[26]: <AxesSubplot:xlabel='value', ylabel='variable'>
```






```
[27]: imports_boxplot_ukfrspde_df = pd.
      ↪ melt(imports_ukfrspde_df[['charges_insurance_and_freight_adj',
      ↪ 'calculated_duties_adj']])
      sns.boxplot(data=imports_boxplot_ukfrspde_df, y='variable', x='value')
```

[27]: <AxesSubplot:xlabel='value', ylabel='variable'>




```
[28]: sns.boxplot(data=imports_ukfrspde_df['First Unit of Quantity'])
```

[28]: <AxesSubplot:>



```
[29]: sns.boxplot(data=imports_ukfrspde_df['wine_quantity_per_capita'])
```

```
[29]: <AxesSubplot:>
```



Let's look at some differences in the duties between the set of countries this additional tariff was put on and the rest of the world.

```
[30]: various_duties_df = imports_world_df[['month']]

various_duties_df['world_average_duties'] =
    (imports_world_df['calculated_duties_adj'] -
    imports_ukfrspde_df['calculated_duties_adj']) /
    (imports_world_df['dutiabale_value_adj'] -
    imports_ukfrspde_df['dutiabale_value_adj'])
various_duties_df['world_average_duties'] =
    various_duties_df['world_average_duties'].astype('f4')

various_duties_df['ukfrspde_average_duties'] =
    (imports_ukfrspde_df['calculated_duties_adj'] -
    imports_ukfrspde_df['dutiabale_value_adj']) /
    (imports_ukfrspde_df['calculated_duties_adj'] -
    imports_ukfrspde_df['dutiabale_value_adj'])
various_duties_df['ukfrspde_average_duties'] =
    various_duties_df['ukfrspde_average_duties'].astype('f4')

duties_plot = sns.lineplot(data=pd.melt(various_duties_df, ['month']),
    x='month', y='value', hue='variable',
    title='Duty Charges for U.S. Wine Imports', xlabel='Month',
    ylabel='Percentage of Dutiable Value')

plt.legend(title='Region', loc='upper left', labels=['World', 'UK', 'FR', 'SP', 'DE'])
plt.show(duties_plot)
```

```
<ipython-input-30-dbdba5509ea0>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
various_duties_df['world_average_duties'] =
(imports_world_df['calculated_duties_adj'] -
imports_ukfrspde_df['calculated_duties_adj']) /
(imports_world_df['dutiabale_value_adj'] -
imports_ukfrspde_df['dutiabale_value_adj'])
```

```
<ipython-input-30-dbdba5509ea0>:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
various_duties_df['world_average_duties'] =
various_duties_df['world_average_duties'].astype('f4')
```

```
<ipython-input-30-dbdba5509ea0>:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
```


Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
various_duties_df['ukfrspde_average_duties'] =  
imports_ukfrspde_df['calculated_duties_adj'] /  
imports_ukfrspde_df['dutiabale_value_adj']  
<ipython-input-30-dbdba5509ea0>:7: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
various_duties_df['ukfrspde_average_duties'] =  
various_duties_df['ukfrspde_average_duties'].astype('f4')
```



```
[31]: various_duties_df['ukfrspde_quantity'] = imports_ukfrspde_df['First Unit of_  
↳Quantity'] / 1000000  
various_duties_df['ukfrspde_quantity'] = various_duties_df['ukfrspde_quantity'].  
↳astype('f4')  
various_duties_df['world_quantity'] = (imports_world_df['First Unit of_  
↳Quantity'] - imports_ukfrspde_df['First Unit of Quantity']) / 1000000
```

```

various_duties_df['world_quantity'] = various_duties_df['world_quantity'].
↳astype('f4')

imports_quantity_plot = sns.lineplot(data=pd.melt(various_duties_df[['month',
↳'world_quantity', 'ukfrspde_quantity']], ['month']), x='month', y='value',
↳hue='variable',)
imports_quantity_plot.set(title='Total Wine Imports to the U.S.',
↳xlabel='Month', ylabel='Millions of Liters')

plt.legend(title='Region', loc='upper left', labels=['World', 'UK, FR, SP, DE'])
plt.show()

```

<ipython-input-31-54b5555526f5>:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

various_duties_df['ukfrspde_quantity'] = imports_ukfrspde_df['First Unit of
Quantity'] / 1000000

```

<ipython-input-31-54b5555526f5>:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

various_duties_df['ukfrspde_quantity'] =
various_duties_df['ukfrspde_quantity'].astype('f4')

```

<ipython-input-31-54b5555526f5>:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

various_duties_df['world_quantity'] = (imports_world_df['First Unit of
Quantity'] - imports_ukfrspde_df['First Unit of Quantity']) / 1000000

```

<ipython-input-31-54b5555526f5>:4: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.


Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

various_duties_df['world_quantity'] =
various_duties_df['world_quantity'].astype('f4')

```



```
[32]: various_duties_df['ukfrspde_quantity'] = imports_ukfrspde_df['First Unit of
↳Quantity']
various_duties_df['ukfrspde_quantity'] = various_duties_df['ukfrspde_quantity'].
↳astype('f4')
various_duties_df['ukfrspde_quantity_per_capita'] =
↳various_duties_df['ukfrspde_quantity'] / imports_ukfrspde_df['population']
various_duties_df['world_quantity'] = (imports_world_df['First Unit of
↳Quantity'] - imports_ukfrspde_df['First Unit of Quantity'])
various_duties_df['world_quantity'] = various_duties_df['world_quantity'].
↳astype('f4')
various_duties_df['world_quantity_per_capita'] =
↳various_duties_df['world_quantity'] / imports_world_df['population']

imports_quantity_plot = sns.lineplot(data=pd.melt(various_duties_df[['month',
↳'world_quantity_per_capita', 'ukfrspde_quantity_per_capita']], ['month']),
↳x='month', y='value', hue='variable',)
imports_quantity_plot.set(title='Wine Imports to the U.S. Per Capita',
↳xlabel='Month', ylabel='Liters per Person')

plt.legend(title='Region', loc='upper left', labels=['World', 'UK, FR, SP, DE'])
plt.show()
```

<ipython-input-32-1b6560be979d>:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
`various_duties_df['ukfrspde_quantity'] = imports_ukfrspde_df['First Unit of Quantity']`  
<ipython-input-32-1b6560be979d>:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead


See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
`various_duties_df['ukfrspde_quantity'] = various_duties_df['ukfrspde_quantity'].astype('f4')`  
<ipython-input-32-1b6560be979d>:3: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
`various_duties_df['ukfrspde_quantity_per_capita'] = various_duties_df['ukfrspde_quantity'] / imports_ukfrspde_df['population']`  
<ipython-input-32-1b6560be979d>:4: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
`various_duties_df['world_quantity'] = (imports_world_df['First Unit of Quantity'] - imports_ukfrspde_df['First Unit of Quantity'])`  
<ipython-input-32-1b6560be979d>:5: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
`various_duties_df['world_quantity'] = various_duties_df['world_quantity'].astype('f4')`  
<ipython-input-32-1b6560be979d>:6: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
`various_duties_df['world_quantity_per_capita'] = various_duties_df['world_quantity'] / imports_world_df['population']`



```
[33]: various_duties_df['ukfrspde_percentage'] = \
    → various_duties_df['ukfrspde_quantity'] / \
    → (various_duties_df['world_quantity'] + \
    → various_duties_df['ukfrspde_quantity'])
various_duties_df['ukfrspde_percentage'] = \
    → various_duties_df['ukfrspde_percentage'].astype('f4')
various_duties_df['percentage_rolling_avg_7'] = various_duties_df.
    → ukfrspde_percentage.rolling(7).mean()

percentage_rolling_avg_plot = sns.lineplot(x='month', \
    → y='percentage_rolling_avg_7', data=various_duties_df)
percentage_rolling_avg_plot.set(title='Percent of UK, Fr, Sp, and De Wine \
    → Imports to U.S. (7-Day Avg)', xlabel='Year', ylabel='Percent of World \
    → Imports')
plt.show()
```

<ipython-input-33-c614e1da5919>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
various\_duties\_df['ukfrspde\_percentage'] =




```
various_duties_df['ukfrspde_quantity'] / (various_duties_df['world_quantity'] +
various_duties_df['ukfrspde_quantity'])
<ipython-input-33-c614e1da5919>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
various_duties_df['ukfrspde_percentage'] =
various_duties_df['ukfrspde_percentage'].astype('f4')
<ipython-input-33-c614e1da5919>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
various_duties_df['percentage_rolling_avg_7'] =
various_duties_df.ukfrspde_percentage.rolling(7).mean()
```



**U.S. Exports** First let's adjust for inflation and population.

```
[34]: exports_df['month'] = exports_df['month'].astype('datetime64[ns]')
exports_df = exports_df.merge(ppi_df, on='month')
```

```

exports_df = exports_df.merge(population_df[['month', 'population']],
    ↪on='month')

exports_df['fas_value_adj'] = exports_df['FAS Value'] /
    ↪exports_df['PCU3121303121300'] * 100
exports_df['quantity_per_capita'] = exports_df['First Unit of Quantity'] /
    ↪exports_df['population']

exports_df.head()

```

```

[34]:      month FAS Value First Unit of Quantity observation_date \
0 2000-01-31 45525687          22385348      2000-01-01
1 2000-02-29 41336082          20954045      2000-02-01
2 2000-03-31 41446856          20240059      2000-03-01
3 2000-04-30 41411758          21745478      2000-04-01
4 2000-05-31 42672115          22059299      2000-05-01

      PCU3121303121300  population  fas_value_adj quantity_per_capita
0              103.0    281083000  44199696.116505          0.07964
1              101.1    281299000  40886332.344214          0.07449
2              103.0    281531000  40239666.019417          0.071893
3              101.0    281763000  41001740.594059          0.077176
4              101.9    281996000  41876462.217861          0.078226

```

```

[35]: display(exports_df.head())
      display(exports_df.describe())
      display(exports_df.info())

```

```

      month FAS Value First Unit of Quantity observation_date \
0 2000-01-31 45525687          22385348      2000-01-01
1 2000-02-29 41336082          20954045      2000-02-01
2 2000-03-31 41446856          20240059      2000-03-01
3 2000-04-30 41411758          21745478      2000-04-01
4 2000-05-31 42672115          22059299      2000-05-01

      PCU3121303121300  population  fas_value_adj quantity_per_capita
0              103.0    281083000  44199696.116505          0.07964
1              101.1    281299000  40886332.344214          0.07449
2              103.0    281531000  40239666.019417          0.071893
3              101.0    281763000  41001740.594059          0.077176
4              101.9    281996000  41876462.217861          0.078226

      PCU3121303121300  population
count          264.000000  2.640000e+02
mean           116.571076  3.093217e+08
std             9.456620  1.525593e+07
min            100.900000  2.810830e+08
25%            105.925000  2.960138e+08

```

```

50%          118.800000  3.109405e+08
75%          123.175000  3.231120e+08
max           134.442000  3.318950e+08

```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 264 entries, 0 to 263
```

```
Data columns (total 8 columns):
```

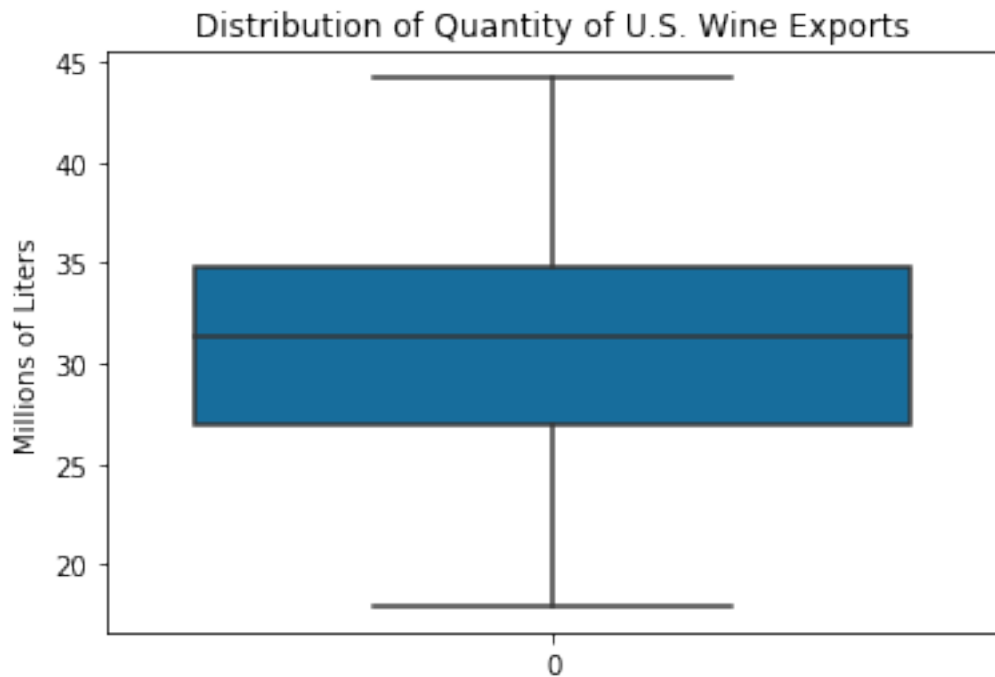
#	Column	Non-Null Count	Dtype
0	month	264 non-null	datetime64[ns]
1	FAS Value	264 non-null	object
2	First Unit of Quantity	264 non-null	object
3	observation_date	264 non-null	datetime64[ns]
4	PCU3121303121300	264 non-null	float64
5	population	264 non-null	int64
6	fas_value_adj	264 non-null	object
7	quantity_per_capita	264 non-null	object

```
dtypes: datetime64[ns](2), float64(1), int64(1), object(4)
```

```
memory usage: 18.6+ KB
```

```
None
```

```
[36]: exports_quantity_boxplot = sns.boxplot(data=exports_df['First Unit of
↳Quantity'] / 1000000)
exports_quantity_boxplot.set(title='Distribution of Quantity of U.S. Wine
↳Exports', ylabel='Millions of Liters')
plt.show()
```




```
[37]: exports_value_boxplot = sns.boxplot(data=exports_df['fas_value_adj'] / 1000000)
exports_value_boxplot.set(title='Distribution U.S. Wine Exports\' Value',
    ↳ylabel='Millions of Dollars')
plt.show()
```




```
[38]: exports_df['first_unit_of_quantity_millions'] = exports_df['First Unit of
    ↳Quantity'].astype('f4') / 1000000
exports_df['quantity_in_mil_rolling_avg_7'] = exports_df.
    ↳first_unit_of_quantity_millions.rolling(7).mean()

exports_quantity_plot = sns.lineplot(data=exports_df, x='month',
    ↳y='quantity_in_mil_rolling_avg_7')
exports_quantity_plot.set(title='U.S. Wine Exports (7-Day Avg)',
    ↳ylabel='Millions of Liters of Wine', xlabel='Year')
plt.show()
```



```
[39]: exports_df['fas_value_millions'] = exports_df['fas_value_adj'].astype('f4') / 1000000
      exports_df['fas_value_millions_rolling_avg_7'] = exports_df.fas_value_millions.rolling(7).mean()

      exports_quantity_plot = sns.lineplot(data=exports_df, x='month', y='fas_value_millions_rolling_avg_7')
      exports_quantity_plot.set(title='U.S. Wine Exports (7-Day Avg)', ylabel='Value of Wine in Millions of Dollars', xlabel='Year')
      plt.show()
```



**Domestic Production** This data is provided by the Alcohol and Tobacco Tax and Trade Bureau (TTB). The data is in gallons instead of liters. Before looking into it, I'm going to quickly convert the timestamp to be the last day of the month and the values to be in liters instead of gallons.

```
[40]: def galons_to_liters(g):
        if math.isnan(g):
            return
        return g * 3.785411784

[41]: production_df['month'] = production_df['month'].astype(str).map(lambda x:
    ↪last_day_of_month(x))
production_df['month'] = production_df['month'].astype('datetime64[ns]')

for c in ['bulk', 'bottled', 'cider', 'effervescent', 'wine_gross']:
    production_df[c] = production_df[c].map(lambda x: galons_to_liters(x))
```

Now let's adjust for per-capita rates.

```
[42]: production_df = production_df.merge(population_df[['month', 'population']],
    ↪on='month')
for c in production_df.columns:
    if c != 'month' and c != 'population':
        col_name = c + '_adj'
        production_df[col_name] = production_df[c] / production_df['population']
```

```
[43]: display(production_df.head())
display(production_df.describe())
display(production_df.info())
```

	month	bulk	bottled	cider	effervescent \
0	2000-01-31	1.131505e+08	1.244070e+08	NaN	6.909175e+06
1	2000-02-29	7.179357e+07	1.375283e+08	NaN	4.377026e+06
2	2000-03-31	4.635628e+07	1.603837e+08	NaN	9.321474e+06
3	2000-04-30	3.296724e+07	1.423004e+08	2.045088e+06	7.881046e+06
4	2000-05-31	3.178035e+07	1.612658e+08	6.959646e+06	6.334834e+06

	wine_gross	population	bulk_adj	bottled_adj	cider_adj \
0	2.444667e+08	281083000	0.402552	0.442599	NaN
1	2.136989e+08	281299000	0.255222	0.488904	NaN
2	2.160614e+08	281531000	0.164658	0.569684	NaN
3	1.811036e+08	281763000	0.117003	0.505036	0.007258
4	1.924214e+08	281996000	0.112698	0.571873	0.024680

	effervescent_adj	wine_gross_adj	population_adj
0	0.024581	0.869731	1.0
1	0.015560	0.759686	1.0
2	0.033110	0.767452	1.0
3	0.027970	0.642752	1.0
4	0.022464	0.682355	1.0

	bulk	bottled	cider	effervescent	wine_gross \
count	2.620000e+02	2.620000e+02	2.590000e+02	2.620000e+02	2.640000e+02
mean	2.366476e+08	1.929920e+08	7.290981e+06	8.266621e+06	4.274359e+08
std	2.329463e+08	3.586958e+07	6.452561e+06	2.698712e+06	2.328960e+08
min	2.661546e+07	1.244070e+08	5.490210e+05	2.678186e+06	0.000000e+00
25%	4.742184e+07	1.640341e+08	1.826624e+06	6.121333e+06	2.391613e+08
50%	1.180365e+08	1.899242e+08	2.811740e+06	8.068083e+06	3.199454e+08
75%	4.498815e+08	2.169504e+08	1.328316e+07	1.011393e+07	6.190754e+08
max	8.200452e+08	3.474637e+08	2.140858e+07	1.481254e+07	1.019108e+09

	population	bulk_adj	bottled_adj	cider_adj	effervescent_adj \
count	2.640000e+02	262.000000	262.000000	259.000000	262.000000
mean	3.093217e+08	0.762839	0.621356	0.022790	0.026546
std	1.525593e+07	0.749701	0.093725	0.019675	0.007954
min	2.810830e+08	0.090648	0.442599	0.001865	0.009267
25%	2.960138e+08	0.156859	0.550870	0.006125	0.020363
50%	3.109405e+08	0.377495	0.611718	0.009091	0.026252
75%	3.231120e+08	1.410859	0.678556	0.040792	0.032392
max	3.318950e+08	2.622399	1.112624	0.066741	0.046178

	wine_gross_adj	population_adj
count	264.000000	264.0
mean	1.377695	1.0

```
std          0.741864          0.0
min          0.000000          1.0
25%          0.782410          1.0
50%          1.034645          1.0
75%          1.999028          1.0
max          3.258976          1.0
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 264 entries, 0 to 263
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	month	264 non-null	datetime64[ns]
1	bulk	262 non-null	float64
2	bottled	262 non-null	float64
3	cider	259 non-null	float64
4	effervescent	262 non-null	float64
5	wine_gross	264 non-null	float64
6	population	264 non-null	int64
7	bulk_adj	262 non-null	float64
8	bottled_adj	262 non-null	float64
9	cider_adj	259 non-null	float64
10	effervescent_adj	262 non-null	float64
11	wine_gross_adj	264 non-null	float64
12	population_adj	264 non-null	float64


```
dtypes: datetime64[ns](1), float64(11), int64(1)
```

```
memory usage: 28.9 KB
```

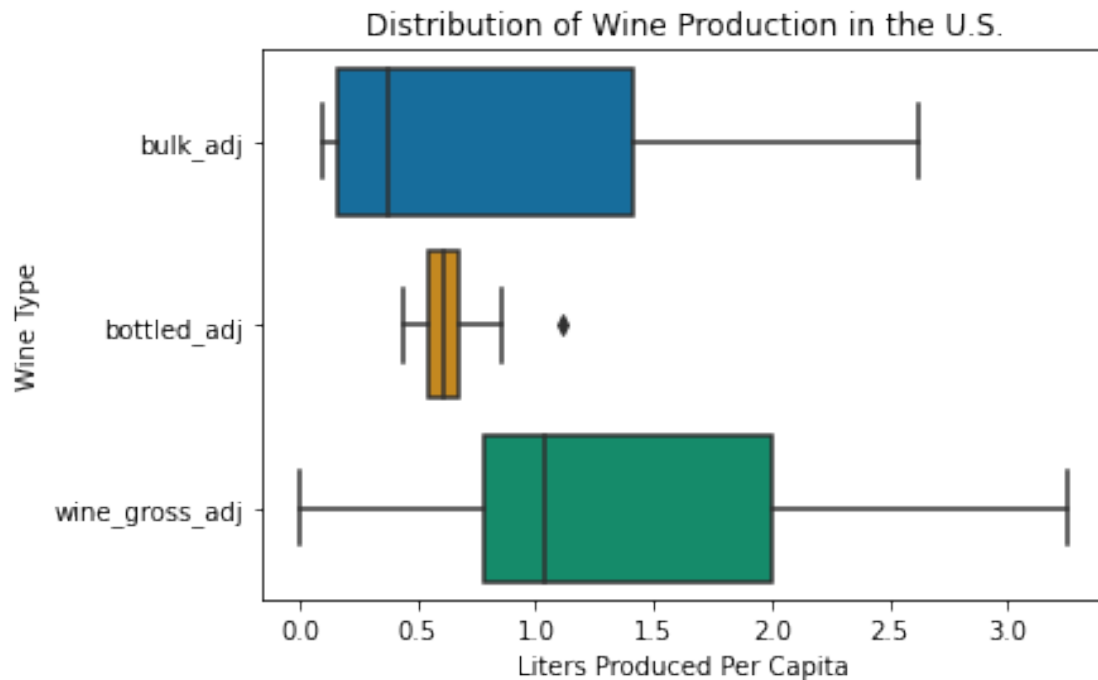
```
None
```

```
[44]: wine_prod_boxplot_df = pd.melt(production_df[['bulk', 'bottled', 'wine_gross']])
wine_prod_boxplot = sns.boxplot(data=wine_prod_boxplot_df, y='variable',
    ↳x='value')
wine_prod_boxplot.set(title='Distribution of Wine Production in the U.S.',
    ↳xlabel='Liters Produced', ylabel='Wine Type')
plt.show()
```

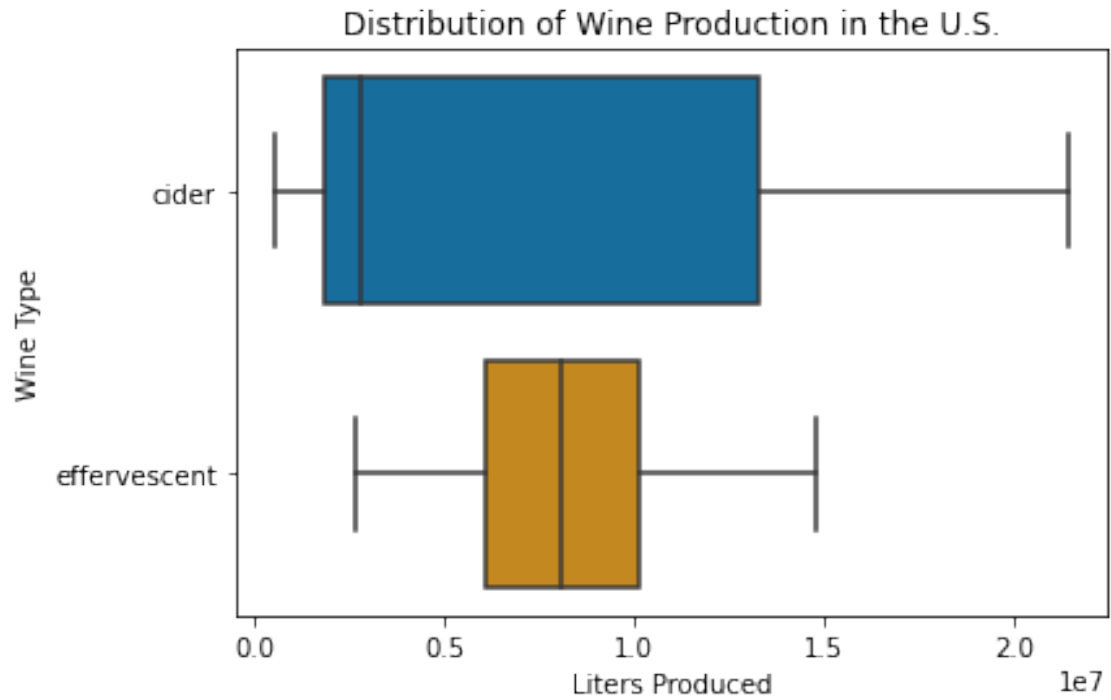




```
[45]: wine_prod_boxplot_adj_df = pd.melt(production_df[['bulk_adj', 'bottled_adj',
    ↳ 'wine_gross_adj']])
wine_prod_boxplot_adj = sns.boxplot(data=wine_prod_boxplot_adj_df,
    ↳ y='variable', x='value')
wine_prod_boxplot_adj.set(title='Distribution of Wine Production in the U.S.',
    ↳ xlabel='Liters Produced Per Capita', ylabel='Wine Type')
plt.show()
```




```
[46]: wine_prod_alt_boxplot_df = pd.melt(production_df[['cider', 'effervescent']])
wine_prod_alt_boxplot = sns.boxplot(data=wine_prod_alt_boxplot_df,
    ↪y='variable', x='value')
wine_prod_alt_boxplot.set(title='Distribution of Wine Production in the U.S.',
    ↪xlabel='Liters Produced', ylabel='Wine Type')
plt.show()
```




```
[47]: production_quantity_plot = sns.lineplot(data=pd.melt(production_df[['month',
    ↳ 'bulk', 'bottled']], ['month']), x='month', y='value', hue='variable',)
production_quantity_plot.set(title='Wine Production in the U.S.',
    ↳ xlabel='Month', ylabel='Liters')

plt.legend(title='Production Type', loc='upper left', labels=['Bulk',
    ↳ 'Bottled'])
plt.show()
```



```
[48]: production_quantity_plot = sns.lineplot(data=pd.melt(production_df[['month',
    ↪ 'cider', 'effervescent']], ['month']), x='month', y='value', hue='variable',)
production_quantity_plot.set(title='Wine Production in the U.S.',
    ↪ xlabel='Month', ylabel='Liters')

plt.legend(title='Alt Wine Type', loc='upper left', labels=['Cider',
    ↪ 'Effervescent'])
plt.show()
```



**Average Wine Prices in U.S. Cities** This data is retrieved from the federal reserve (FRED).

```
[49]: # Standardize dates
wine_prices_df['month'] = wine_prices_df['month'].map(lambda x:
↳last_day_of_month(x))
wine_prices_df['month'] = wine_prices_df['month'].astype('datetime64[ns]')

# Adjust for inflation
wine_prices_df = wine_prices_df.merge(ppi_df, on='month')
wine_prices_df['price_adj'] = wine_prices_df['price'] /
↳exports_df['PCU3121303121300'] * 100
```

```
[50]: display(wine_prices_df.head())
display(wine_prices_df.describe())
display(wine_prices_df.info())
```

	month	price	observation_date	PCU3121303121300	price_adj
0	2000-01-31	5.458	2000-01-01	103.0	5.299029
1	2000-02-29	5.256	2000-02-01	101.1	5.198813
2	2000-03-31	5.471	2000-03-01	103.0	5.311650
3	2000-04-30	5.156	2000-04-01	101.0	5.104950
4	2000-05-31	5.530	2000-05-01	101.9	5.426889

price PCU3121303121300 price\_adj

count	264.000000	265.000000	264.000000
mean	9.740682	116.637834	8.241252
std	2.578523	9.501049	1.652331
min	5.156000	100.900000	5.104950
25%	7.386500	106.000000	6.774410
50%	10.025000	118.800000	8.631385
75%	12.197750	123.400000	9.803208
max	14.420000	134.442000	11.848809

<class 'pandas.core.frame.DataFrame'>

Int64Index: 265 entries, 0 to 264

Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	month	265 non-null	datetime64[ns]
1	price	264 non-null	float64
2	observation_date	265 non-null	datetime64[ns]
3	PCU3121303121300	265 non-null	float64
4	price_adj	264 non-null	float64

dtypes: datetime64[ns](2), float64(3)


memory usage: 20.5 KB

None

```
[51]: wine_prices_df['prices_rolling_avg_7'] = wine_prices_df.price_adj.rolling(7).
      ↪mean()
```

```
wine_prices_plot = sns.lineplot(data=wine_prices_df, x='month',
      ↪y='prices_rolling_avg_7')
wine_prices_plot.set(title='U.S. Wine Prices (7-Day Avg)', ylabel='Dollars',
      ↪xlabel='Year')
```

```
[51]: [Text(0.5, 1.0, 'U.S. Wine Prices (7-Day Avg)'),
      Text(0, 0.5, 'Dollars'),
      Text(0.5, 0, 'Year')]
```



## 1.4 Combine Datasets

Construct a master dataset.

```
[63]: tariff_df.head()
```

```
[63]:      month  frspger_25
0  2000-01         0
1  2000-02         0
2  2000-03         0
3  2000-04         0
4  2000-05         0
```

```
[66]: master_df = pd.DataFrame()
master_df['month'] = months

# Population
master_df = master_df.merge(population_df[['month', 'population']], on='month')

# Prices
master_df = master_df.merge(wine_prices_df[['month', 'price', 'price_adj']],
    on='month')

# Domestic Production
```

```

master_df = master_df.merge(production_df[['month', 'bulk', 'bottled', 'cider',
↳ 'effervescent', 'wine_gross', 'bulk_adj', 'bottled_adj', 'cider_adj',
↳ 'effervescent_adj', 'wine_gross_adj']], on='month')

# Exports
## renaming columns
exports_df.rename(columns={
    'First Unit of Quantity': 'quantity_exports',
    'fas_value_adj': 'fas_value_adj_exports'
}, inplace=True)
## merging
master_df = master_df.merge(exports_df[['month', 'quantity_exports',
↳ 'fas_value_adj_exports']], on='month')

# Imports
## renaming world imports columns
imports_world_df.rename(columns={
    'Dutiable Value': 'dutiable_value_world_imports',
    'dutiable_value_adj': 'dutiable_value_adj_world_imports',
    'Landed Duty-Paid Value': 'landed_duty_paid_value_world_imports',
    'landed_duty_paid_value_adj': 'landed_duty_paid_value_adj_world_imports',
    'Customs Value': 'customs_value_world_imports',
    'customs_value_adj': 'customs_value_adj_world_imports',
    'First Unit of Quantity': 'quantity_world_imports',
    'Charges, Insurance, and Freight':
↳ 'charges_insurance_freight_world_imports',
    'charges_insurance_and_freight_adj':
↳ 'charges_insurance_freight_adj_world_imports',
    'Calculated Duties': 'calculated_duties_world_imports',
    'calculated_duties_adj': 'calculated_duties_adj_world_imports'
}, inplace=True)
## renaming ukfspde imports columns
imports_ukfspde_df.rename(columns={
    'Dutiable Value': 'dutiable_value_ukfspde_imports',
    'dutiable_value_adj': 'dutiable_value_adj_ukfspde_imports',
    'Landed Duty-Paid Value': 'landed_duty_paid_value_ukfspde_imports',
    'landed_duty_paid_value_adj': 'landed_duty_paid_value_adj_ukfspde_imports',
    'Customs Value': 'customs_value_ukfspde_imports',
    'customs_value_adj': 'customs_value_adj_ukfspde_imports',
    'First Unit of Quantity': 'quantity_ukfspde_imports',
    'Charges, Insurance, and Freight':
↳ 'charges_insurance_freight_ukfspde_imports',
    'charges_insurance_and_freight_adj':
↳ 'charges_insurance_freight_adj_ukfspde_imports',
    'Calculated Duties': 'calculated_duties_ukfspde_imports',
    'calculated_duties_adj': 'calculated_duties_adj_ukfspde_imports'
})

```



```

    }, inplace=True)
## merging
master_df = master_df.merge(imports_world_df[['month',
    ↳ 'dutiable_value_world_imports',
    ↳ 'dutiable_value_adj_world_imports', 'landed_duty_paid_value_world_imports',
    ↳ 'landed_duty_paid_value_adj_world_imports', 'customs_value_world_imports',
    ↳ 'customs_value_adj_world_imports', 'quantity_world_imports',
    ↳ 'charges_insurance_freight_world_imports',
    ↳ 'charges_insurance_freight_adj_world_imports',
    ↳ 'calculated_duties_world_imports', 'calculated_duties_adj_world_imports']],
    ↳ on='month')
master_df = master_df.merge(imports_ukfrspde_df[['month',
    ↳ 'dutiable_value_ukfrspde_imports',
    ↳ 'dutiable_value_adj_ukfrspde_imports',
    ↳ 'landed_duty_paid_value_ukfrspde_imports',
    ↳ 'landed_duty_paid_value_adj_ukfrspde_imports',
    ↳ 'customs_value_ukfrspde_imports',
    ↳ 'customs_value_adj_ukfrspde_imports', 'quantity_ukfrspde_imports',
    ↳ 'charges_insurance_freight_ukfrspde_imports',
    ↳ 'charges_insurance_freight_adj_ukfrspde_imports',
    ↳ 'calculated_duties_ukfrspde_imports',
    ↳ 'calculated_duties_adj_ukfrspde_imports']], on='month')

# Tariff Enacted Boolean
tariff_df['month'] = tariff_df['month'].map(lambda x: last_day_of_month(str(x)))
tariff_df['month'] = tariff_df['month'].astype('datetime64[ns]')
master_df = master_df.merge(tariff_df, on='month')

```

```

[66]: Index(['month', 'population', 'price', 'price_adj', 'bulk', 'bottled', 'cider',
    'effervescent', 'wine_gross', 'bulk_adj', 'bottled_adj', 'cider_adj',
    'effervescent_adj', 'wine_gross_adj', 'quantity_exports',
    'fas_value_adj_exports', 'dutiable_value_world_imports',
    'dutiable_value_adj_world_imports',
    'landed_duty_paid_value_world_imports',
    'landed_duty_paid_value_adj_world_imports',
    'customs_value_world_imports', 'customs_value_adj_world_imports',
    'quantity_world_imports', 'charges_insurance_freight_world_imports',
    'charges_insurance_freight_adj_world_imports',
    'calculated_duties_world_imports',
    'calculated_duties_adj_world_imports',
    'dutiable_value_ukfrspde_imports',
    'dutiable_value_adj_ukfrspde_imports',
    'landed_duty_paid_value_ukfrspde_imports',
    'landed_duty_paid_value_adj_ukfrspde_imports',
    'customs_value_ukfrspde_imports', 'customs_value_adj_ukfrspde_imports',
    'quantity_ukfrspde_imports',

```

```
'charges_insurance_freight_ukfrspde_imports',
'charges_insurance_freight_adj_ukfrspde_imports',
'calculated_duties_ukfrspde_imports',
'calculated_duties_adj_ukfrspde_imports', 'frspger_25'],
dtype='object')
```

The dataset looks pretty good. I'll export this as an excel file just to hold on to it as is.

```
[67]: display(master_df.head())
display(master_df.describe())
display(master_df.info())
```

	month	population	price	price_adj	bulk	bottled	\
0	2000-01-31	281083000	5.458	5.299029	1.131505e+08	1.244070e+08	
1	2000-02-29	281299000	5.256	5.198813	7.179357e+07	1.375283e+08	
2	2000-03-31	281531000	5.471	5.311650	4.635628e+07	1.603837e+08	
3	2000-04-30	281763000	5.156	5.104950	3.296724e+07	1.423004e+08	
4	2000-05-31	281996000	5.530	5.426889	3.178035e+07	1.612658e+08	

	cider	effervescent	wine_gross	bulk_adj	...	\
0	NaN	6.909175e+06	2.444667e+08	0.402552	...	
1	NaN	4.377026e+06	2.136989e+08	0.255222	...	
2	NaN	9.321474e+06	2.160614e+08	0.164658	...	
3	2.045088e+06	7.881046e+06	1.811036e+08	0.117003	...	
4	6.959646e+06	6.334834e+06	1.924214e+08	0.112698	...	

	landed_duty_paid_value_ukfrspde_imports	\
0	59812261	
1	77087577	
2	87219165	
3	87040067	
4	79534639	

	landed_duty_paid_value_adj_ukfrspde_imports	\
0	58070156.31068	
1	76248839.762611	
2	84678800.970874	
3	86178284.158416	
4	78051657.50736	

	customs_value_ukfrspde_imports	customs_value_adj_ukfrspde_imports	\
0	56706027	55054395.145631	
1	73873201	73069437.1909	
2	83500840	81068776.699029	
3	83075769	82253236.633663	
4	75599523	74189914.622179	

	quantity_ukfrspde_imports	charges_insurance_freight_ukfrspde_imports	\
--	---------------------------	--	---

0	8768676	2406931
1	8961916	2356486
2	10474993	2804317
3	11128077	2989501
4	10874051	3037785

charges_insurance_freight_adj_ukfrspde_imports \	
0	2336826.213592
1	2330846.686449
2	2722637.864078
3	2959901.980198
4	2981143.277723

calculated_duties_ukfrspde_imports calculated_duties_adj_ukfrspde_imports \		
0	699303	678934.951456
1	857890	848555.885262
2	914008	887386.407767
3	974797	965145.544554
4	897331	880599.607458

frspger_25	
0	0
1	0
2	0
3	0
4	0

[5 rows x 39 columns]

	population	price	price_adj	bulk	bottled \
count	2.640000e+02	264.000000	264.000000	2.620000e+02	2.620000e+02
mean	3.093217e+08	9.740682	8.241252	2.366476e+08	1.929920e+08
std	1.525593e+07	2.578523	1.652331	2.329463e+08	3.586958e+07
min	2.810830e+08	5.156000	5.104950	2.661546e+07	1.244070e+08
25%	2.960138e+08	7.386500	6.774410	4.742184e+07	1.640341e+08
50%	3.109405e+08	10.025000	8.631385	1.180365e+08	1.899242e+08
75%	3.231120e+08	12.197750	9.803208	4.498815e+08	2.169504e+08
max	3.318950e+08	14.420000	11.848809	8.200452e+08	3.474637e+08

	cider	effervescent	wine_gross	bulk_adj	bottled_adj \
count	2.590000e+02	2.620000e+02	2.640000e+02	262.000000	262.000000
mean	7.290981e+06	8.266621e+06	4.274359e+08	0.762839	0.621356
std	6.452561e+06	2.698712e+06	2.328960e+08	0.749701	0.093725
min	5.490210e+05	2.678186e+06	0.000000e+00	0.090648	0.442599
25%	1.826624e+06	6.121333e+06	2.391613e+08	0.156859	0.550870
50%	2.811740e+06	8.068083e+06	3.199454e+08	0.377495	0.611718
75%	1.328316e+07	1.011393e+07	6.190754e+08	1.410859	0.678556
max	2.140858e+07	1.481254e+07	1.019108e+09	2.622399	1.112624

	cider_adj	effervescent_adj	wine_gross_adj	frspger_25
count	259.000000	262.000000	264.000000	264.000000
mean	0.022790	0.026546	1.377695	0.079545
std	0.019675	0.007954	0.741864	0.271102
min	0.001865	0.009267	0.000000	0.000000
25%	0.006125	0.020363	0.782410	0.000000
50%	0.009091	0.026252	1.034645	0.000000
75%	0.040792	0.032392	1.999028	0.000000
max	0.066741	0.046178	3.258976	1.000000

<class 'pandas.core.frame.DataFrame'>

Int64Index: 264 entries, 0 to 263

Data columns (total 39 columns):

#	Column	Non-Null Count	Dtype
0	month	264 non-null	
	datetime64[ns]		
1	population	264 non-null	int64
2	price	264 non-null	float64
3	price_adj	264 non-null	float64
4	bulk	262 non-null	float64
5	bottled	262 non-null	float64
6	cider	259 non-null	float64
7	effervescent	262 non-null	float64
8	wine_gross	264 non-null	float64
9	bulk_adj	262 non-null	float64
10	bottled_adj	262 non-null	float64
11	cider_adj	259 non-null	float64
12	effervescent_adj	262 non-null	float64
13	wine_gross_adj	264 non-null	float64
14	quantity_exports	264 non-null	object
15	fas_value_adj_exports	264 non-null	object
16	dutiable_value_world_imports	264 non-null	object
17	dutiable_value_adj_world_imports	264 non-null	object
18	landed_duty_paid_value_world_imports	264 non-null	object
19	landed_duty_paid_value_adj_world_imports	264 non-null	object
20	customs_value_world_imports	264 non-null	object
21	customs_value_adj_world_imports	264 non-null	object
22	quantity_world_imports	264 non-null	object
23	charges_insurance_freight_world_imports	264 non-null	object
24	charges_insurance_freight_adj_world_imports	264 non-null	object
25	calculated_duties_world_imports	264 non-null	object
26	calculated_duties_adj_world_imports	264 non-null	object
27	dutiable_value_ukfrspde_imports	264 non-null	object
28	dutiable_value_adj_ukfrspde_imports	264 non-null	object
29	landed_duty_paid_value_ukfrspde_imports	264 non-null	object
30	landed_duty_paid_value_adj_ukfrspde_imports	264 non-null	object

```

31  customs_value_ukfrspde_imports          264 non-null  object
32  customs_value_adj_ukfrspde_imports      264 non-null  object
33  quantity_ukfrspde_imports               264 non-null  object
34  charges_insurance_freight_ukfrspde_imports 264 non-null  object
35  charges_insurance_freight_adj_ukfrspde_imports 264 non-null  object
36  calculated_duties_ukfrspde_imports       264 non-null  object
37  calculated_duties_adj_ukfrspde_imports    264 non-null  object
38  frspger_25                              264 non-null  int64
dtypes: datetime64[ns](1), float64(12), int64(2), object(24)
memory usage: 82.5+ KB

None

```

```
[68]: master_df.to_excel('../data/master-data.xlsx')
```